Network modeling – AY 2026/2027
Homework 1 – Modeling Zoom video uplink traffic

Alessandro Vulcu – ID 2185190
University of Padua

Abstract

I analyze one outbound UDP flow from a Zoom meeting between a PC and an iPad. Packets occur in tight clusters, which a Poisson source fails to reproduce. An empirical batch-renewal model captures packet-gap and size distributions and matches the mean rate on a held-out segment. It does not reproduce the observed correlation of traffic over longer time scales.

1 Measurement and sample selection

I captured Zoom traffic with Wireshark on the PC's Ethernet interface `enp5s0`; the iPad was a second participant on the same LAN. The experiment progressed through idle/muted operation, a connected meeting with media off, microphone-only speech and silence, camera-on speech, and camera off again. These phase durations were approximate: the stopwatch ran for about ten minutes, so I do not treat the planned durations as exact boundaries.
The video-related flow was identified by switching the PC camera off and on and observing its traffic decrease and recover. Its endpoints were `192.168.3.7:42613` and `206.247.58.75:8801`. I selected only the PC-to-server direction using

```
udp.stream == 0 && ip.src == 192.168.3.7 && frame.number > 32864.
```

The filter retains 34,319 packets spanning 615.08 s. I set $t = 0$ at the first selected packet (431.2223 s into the raw capture), not at the stopwatch start. The rate rises near $t = 260$ s and falls near $t = 578$ s (Fig. 1). I therefore select the interior camera-on interval $[270, 570)$ s: 32,136 packets, 300 s, and 862.50 kbit/s. Its first 180 s are used for estimation and its last 120 s for validation; these segments are never pooled during fitting. No selected packet is truncated. Packet sizes are `frame.len` in bytes, including captured Ethernet headers, not encoded media payload.

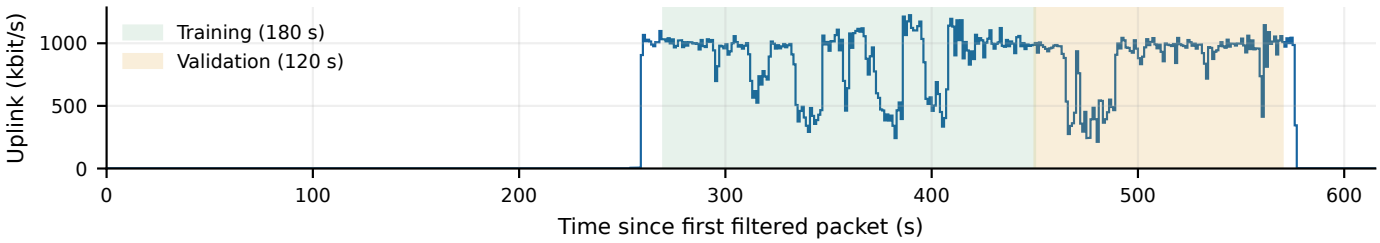


Figure 1: One-second uplink rate and the selected training and validation intervals. Camera transitions are inferred from traffic; audio-only phase boundaries are not resolved in this flow.

2 Source model

In the training segment, packet gaps have mean 9.162 ms, median 5.38 μs, and coefficient of variation 1.711; 63.17% are at most 1 ms. I first fit a Poisson baseline: $A_i \sim \mathrm{Exp}(109.145\,\mathrm{s}^{-1})$, $B_i \equiv 1$, and independent packet sizes sampled from the training empirical distribution. Its exponential gaps cannot represent the concentration near zero.
Following Lessons 3–4, I model the bursty VBR source with a marked and weighted renewal process, augmented with packet offsets. Consecutive packets separated by at most $\delta = 1$ ms belong to the same cluster. A cluster is an observation rule, not an identified codec frame. Let τ_i be its first timestamp, B_i its packet count, $U_{i,k}$ its packet offsets, and $W_{i,k} = 8L_{i,k}$ its weights in bits. Each complete training cycle stores the joint tuple

$$C_i = (D_i, B_i, \{U_{i,k}, W_{i,k}\}_{k=1}^{B_i}), \qquad D_i = \tau_{i+1} - \tau_i.$$

I exclude the boundary clusters, leaving 7,233 cycles. To generate a trace, I draw cycles independently and uniformly from this empirical distribution, emit packets at $\tau_i + U_{i,k}$, then advance by D_i. Thus batch inter-arrivals satisfy $A_{i+1} = D_i$; weights, counts and offsets remain dependent within a sampled cycle. Offsets prevent artificial zero packet gaps. The fitted means are $E[D] = 24.881$ ms and $E[B] = 2.716$ packets, with predicted rate $E[\sum_k W_{i,k}]/E[D] = 865.95$ kbit/s. Thresholds of 0.5 and 3 ms give mean batch sizes 2.707 and 2.774: the clustering is reasonably stable locally.

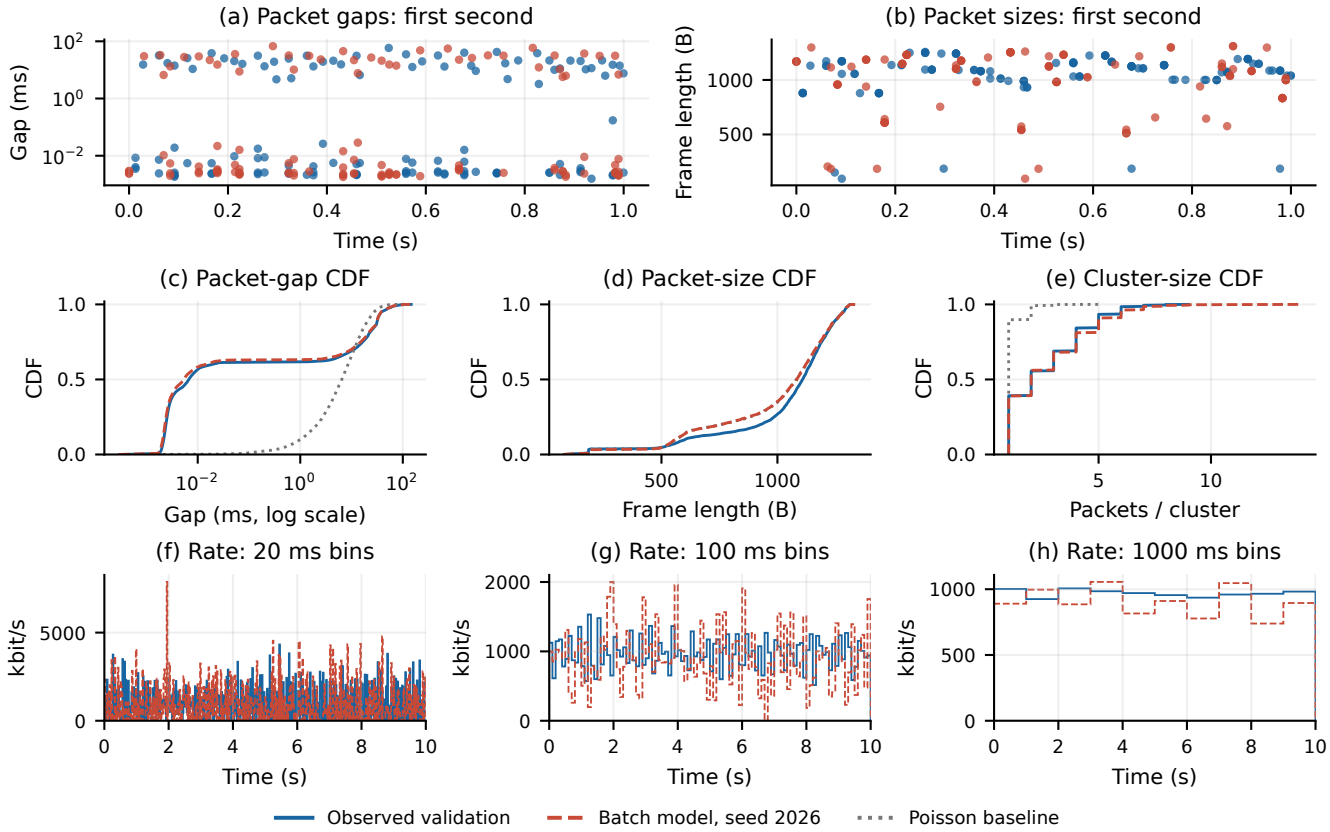


Figure 2: Held-out observations and independent 120 s simulations (seed 2026). Scatterplots show the first second; CDFs use the full traces; rate plots show the first ten seconds. Cluster sizes in (e) use the same 1 ms rule for every trace. Synthetic events are not expected to align with observed events.

3 Results

120 s validation comparison	Observed	Poisson	Batch model
Packet count	12,491	12,927	13,102
Mean frame length (B)	1,029.57	989.97	990.38
Mean rate (kbit/s)	857.35	853.16	865.06
Packet-gap CDF distance	—	0.611	0.050
Count Fano factor, 100 ms bins	1.115	0.995	2.475
Lag-1 count correlation, 1 s bins	0.684	0.151	0.049

Both simulations use training data only and stop at 120 s; packet counts are stochastic, not forced to match observations. The batch realization's mean rate is 0.90% above validation. Across 30 seeds (2026–2055), its mean rate is 864.30 kbit/s; the empirical 5th–95th percentile range is 847.42–879.97 kbit/s. Figure 2 shows that the batch model preserves short packet gaps much better than Poisson: the maximum absolute empirical CDF difference falls from 0.611 to 0.050. This is a descriptive distance, not an independence-based significance test. The size-CDF distance is 0.089 for the batch model, reflecting some shift in sizes between training and validation.

4 Discussion

The batch model captures local clustering and average load, but overestimates 100 ms count variability and loses the strong correlation between successive one-second counts. Independent cycle sampling discards ordering across clusters; it therefore misses slow activity or adaptation changes. A state-dependent or block model would be a next refinement, not a result established here. This single run covers one video-related UDP flow, not all Zoom traffic or separately identified audio. Phase boundaries are inferred and payloads are encrypted. The model describes packets observed at the host capture interface, not the encoder's internal generation times or a universal Zoom traffic law. Controlled camera changes alone do not provide enough cycles to estimate an On–Off sojourn-time model. Course basis: Lesson 3, *Network abstraction and traffic source models*; Lesson 4, *Traffic source models*.